

The Triadic Measurement Substrate

$$\bar{B} = 3\bar{\alpha}$$

Volume 1

Paper 1: Emergent 3D Information Topology and the Binary Alphabet via [[4,2,2]] Quantum Error Detection Structural Correspondence

Aaron Switzer

2026

*Excerpted from the complete two-volume corpus, The Triadic Measurement Substrate. Full corpus and reproducibility
package: <https://doi.org/10.5281/zenodo.22035422>*

Paper 1

Emergent 3D Information Topology and the Binary Alphabet via $[[4,2,2]]$ Quantum Error Detection Structural Correspondence

Aaron Switzer

May 2026

Keywords: Information Theory, FCC Lattice, $[[4,2,2]]$ Code, Geometric Necessity, Causal Succession, Planck Length, Binary Alphabet, Generalized Uncertainty Principle, Computational Irreducibility

Status of Results

This paper distinguishes three categories of claim, each labeled where it appears in the text.

Derived. Results that follow from the five axioms and the Principle of Generative Economy by explicit construction or proof. The triadic minimum $\bar{B} = 3\bar{\alpha}_2$ (§ 2.1), the hexagonal packing of $M\alpha$ (§ 2.2), FCC necessity over HCP via the causal distinguishability argument (§ 2.4), the $1 \rightarrow 4$ multiplier and its remainder-free integer division (§ 4.2), the $[[4,2,2]]$ structural correspondence and code distance 2 (§ V), the wave equation $\partial^2 \varphi / \partial t^2 - c^2 \nabla^2 \varphi = 0$ in the continuum limit (§ VI), and the propagation speed $c_{\text{TMS}} = 1/\sqrt{2}$ from FCC face-diagonal geometry (§ 6.5) are derived in this sense. The scale-invariant $k = 3n$ stabilizer null-space law (§ 5.4) is derived here as a structural property of FCC + $[[4,2,2]]$ dynamics, empirically verified through Paper 5 Appendix E at levels 1 and 2. The sublattice bifurcation corollary (§ 2.4) is derived here in skeleton form, with the full geometric development and empirical verification in Paper 5 § 6.1.3 and Appendix E. The conjugate pair (φ, Π) and the second-order dynamics are derived from signal conservation under FCC connectivity (§ 6.3), with the uniqueness of the second equation tightened in that section.

Structural correspondence. Results in which a TMS construction faithfully reproduces the skeleton of a known framework without claiming reproduction of every feature of that framework. The $[[4,2,2]]$ mapping of § V is the canonical example: the stabilizer algebra, the two logical degrees of freedom, and the code distance are reproduced exactly; the full quantum amplitude formalism is deferred to Paper 2 and subsequent work. The identification $\ell_0 \equiv \ell_p$ in § 3.2 is a structural correspondence in the same sense — the minimum resolvable length of the emergent geometry is identified with the Planck length, with the derivation of $\hbar$ and G from the confirmation field dynamics deferred (§ 3.2 Remark).

Predictions / Conjectures. Falsifiable claims that follow from the framework but whose closed forms are not derived within this paper. The GUP scaling $\beta = f(d, K_\psi)$ (§ 7.2), the computational irreducibility connection of § 7.2a, and the unified geometric account of decoherence (§ 7.3) are predictions in this sense. They are consequences of the framework’s geometric architecture; they are not derivations of f , of explicit decoherence rates, or of falsification thresholds. The functional form of f is refined across Papers 2–4 and pinned to numerical content in Paper 5’s simulation work.

This three-way labeling applies throughout the volume. Every claim in the body text falls into exactly one category, and every category's status is visible at the point of claim.

Notation

Throughout this paper, the bar notation distinguishes confirmed (locked) elements from latent (unlocked) elements. $\bar{\alpha}$ denotes an agent participating in a confirmation event; $\bar{B}$ denotes a confirmed (actualized) bit. Unbarred α and B denote their latent counterparts. Dimensional subscripts $\bar{\alpha}_2$ and $\bar{\alpha}_3$ distinguish 2D and 3D confirmation contexts respectively. A confirmed bit carries a binary value $\bar{B} \in \{0, 1\}$; its signed representation $\chi \equiv 2\bar{B} - 1 \in \{-1, +1\}$ is referred to as its polarization.

Abstract. We propose the Triadic Measurement Substrate (TMS), a generative model in which wave dynamics and a binary computational alphabet emerge necessarily from the interaction between a Maximum Entropy Stochastic Field (U_{sh}) and a 2D hexagonal agent manifold ($M\alpha$). Beginning from five axioms and the Principle of Generative Economy, we demonstrate that hexagonal close-packing, the $1 \rightarrow 4$ informational multiplier, FCC lattice geometry, the binary alphabet $\bar{B} \in \{0, 1\}$ of confirmed bit values, the conjugate field structure, and the $[[4,2,2]]$ quantum error-detection structural correspondence follow by geometric necessity. The fundamental agent spacing is identified with the Planck length ℓ_p , the flop time with the Planck time t_p , and the lattice propagation speed $c_{\text{TMS}} = 1/\sqrt{2}$ is forced by the geometry of FCC face-diagonal connectivity. The continuum limit reproduces the standard wave equation with speed c . The $[[4,2,2]]$ structural correspondence is shown to faithfully encode both the value of every confirmed bit and the inheritance of that value through the causal stack — two independent logical degrees of freedom jointly protected by the two stabilizer generators against any single-agent disruption from U_{sh} . The TMS contains no free dimensionless numerical parameters — every dimensionless quantity is fixed by the axioms and the Principle of Generative Economy — and takes a single dimensionful input (one action scale, $\hbar$), from which the speed of light and Newton's constant follow (§ 3.2–3.3). The absolute length scale is fixed by identifying the substrate's minimum length with the Planck length. The wave equation is not fitted to the model — it is the model's necessary large-scale behavior. The paper closes with a novel falsifiable prediction: the GUP deformation parameter β is not a universal constant but scales with the causal depth and polarization complexity of the confirmation record, providing a geometric and computational derivation of both the Generalized Uncertainty Principle and quantum decoherence from a single structural property of the FCC causal stack.

This paper constitutes Volume 1, Paper 1 of a larger work. Subsequent papers in Volume 1 develop the computational dynamics of the binary alphabet, the emergence of learning and hierarchy, subsystem development, and simulation. Volume 2 addresses the agent-side complement.

I. Definitions and Ontological Primitives

The Principle of Generative Economy

All structure selects the minimum sufficient configuration that maximizes generative potential. This principle governs all selection throughout the TMS — of primitives, confirmation geometry, and lattice structure.

This principle belongs to the deep tradition in physics of parsimony-as-law. Hamilton's Principle of Least Action (Hamilton, 1834) holds that nature selects the path that extremizes the action integral. Jaynes' Maximum Entropy Principle (Jaynes, 1957) holds that unbiased inference selects the least-structured distribution consistent with known constraints. Occam's Razor, in its formal instantiation as Minimum Description Length (Rissanen, 1978), holds that the simplest model consistent with the data is preferred. These principles share a common structure: nature does not take the long way home.

The Principle of Generative Economy is a specific instantiation of this tradition, distinguished by a dual constraint. Existing parsimony principles minimize in one direction — cost, assumptions, description length. The Principle of Generative Economy jointly minimizes structural complexity and maximizes generative capacity. This dual constraint is necessary because pure minimization selects the degenerate case: a point is simpler than a sphere, but generates nothing; a line is simpler than a triangle, but closes nothing. The sphere and the triangle are not selected because they are small — they are selected because they are the minimum configurations that remain generatively non-trivial. Kauffman’s adjacent possible (Kauffman, 1993) is the closest conceptual relative: the idea that systems naturally explore the maximum frontier of what their current structure makes accessible. A parallel framing of dual parsimony has been developed independently by Ma, Tsao, and Shum (2022) for the emergence of intelligence, where Parsimony and Self-Consistency together govern what to learn and how to learn. The Principle of Generative Economy is the geometric form that the parsimony tradition takes when the only primitive is the sphere and the only operation is confirmation. In giving condition (i) below a specific computable form — the moduli dimension of a configuration under axiom-forced constraints — the framework instantiates this established tradition rather than inventing a new measure of simplicity: the reading of complexity as a count of free parameters, and of parsimony as its minimization, is standard across the description-length and model-selection literature cited above. What is particular to this work is not the measure but its application, and specifically the requirement that the constraint set be derived from the axioms, developed below.

Formal Statement (Principle of Generative Economy). Let K denote the set of constraints imposed on a given selection by the axioms and by prior structure relevant to that selection. Let $\mathcal{C}_K$ denote the set of configurations satisfying K . A configuration $C^* \in \mathcal{C}_K$ is the PGE-selection if and only if:

(i) *Minimum specification.* C^* minimizes the *specification cost* over $\mathcal{C}_K$, where the specification cost of a configuration is the number of free real parameters required to fix it beyond what K forces — that is, the dimension of the space of K -satisfying configurations equivalent to it under the symmetry K permits (its moduli dimension). For discrete selections this reduces to an integer count of independent structural commitments; for continuous selections it is the dimension of the orientation-and-shape moduli. The measure is standard and computable: for example, in the selection of the geometric primitive the sphere has specification cost zero (isotropic: no orientation or shape parameter), while a cube costs three, an ellipsoid five, and an irregular form more — so the sphere is the unique minimizer.

(ii) *Generative non-triviality.* C^* admits downstream structure under the substrate dynamics. It is not a generative dead-end.

When K admits a unique configuration satisfying both (i) and (ii), that configuration is the unique PGE-selection. When (i) admits multiple candidates, (ii) is the tiebreaker. When the unique minimum-specification candidate fails (ii), K is too restrictive for generative non-triviality and the axioms must supply further constraint.

The constraint-set requirement. The entire evidential content of a PGE-selection lies not in the parameter count, which is unambiguous, but in the specification of K . Accordingly this work adopts a strict requirement: at every invocation, K must be *derived from the axioms and declared before the selection*, never reverse-engineered from the desired configuration. A PGE-selection is rigorous exactly to the extent that its K is exhibited as axiom-forced. As a worked demonstration on a contested case, the selection of FCC over hexagonal close-packing (§ 2.4) uses $K = \{\text{Axiom 4 (no overwrite) together with the definition of a confirmation datum as a state individuated by its causal relations}\}$; the causal-distinguishability constraint that forbids the ABAB stacking is not an added assumption but a consequence of Axiom 4, since two causally indistinguishable layers are, by the definition of the record, the same recorded state — which Axiom 4 forbids.

Establishing the axiom-forced character of K at each remaining invocation, grouped by constraint-type, is the principal open formal task of the program; the selection functional itself is fixed and computable.

Remark 0.0.1 (On the Two Conditions). Condition (i) is the substantive content of *minimum sufficient*: sufficiency is set by K , minimality is the moduli dimension not forced by K . Condition (ii) is the substantive content of *maximizes generative potential*: generatively sterile candidates are excluded even when minimal. The pair is necessary because pure minimization selects the degenerate case as noted above. Every subsequent invocation of “by Generative Economy” in this volume is to be verified against this formal statement — the constraint set K must be specifiable and axiom-derived, and C^* must satisfy both (i) and (ii) uniquely within $\mathcal{C}_K$.

Illustration (The Sphere Primitive). As a first application: for the selection of the geometric primitive in § 1.2, $K = \{\text{isotropy, geometric primitive, supports the measurement role required by Axiom 1}\}$. Within $\mathcal{C}_K$, the sphere has zero orientation parameters; any other candidate (cube, ellipsoid, regular polyhedron, irregular form) introduces at least one orientation or shape parameter not forced by K . The sphere therefore satisfies (i) uniquely. It satisfies (ii) — it encloses interstices, supports contact geometry, packs in 2D and 3D, and admits the cascade dynamics of subsequent sections. The sphere is the unique PGE-selection. Axiom 2 records this result.

Axiom 0 — Information Fundamentalism

Information is fundamental. It exists as potential prior to confirmation. This potential field is U_{sh} — pure informational superposition, algorithmically incompressible, with complexity $K(n) \approx n$ (Kolmogorov, 1965). It is maximum entropy by definition: a unique state, not a tunable quantity. U_{sh} has no free parameters. This axiom is the formal expression of the informational priority articulated by Wheeler’s “it from bit” (Wheeler, 1990): the physical arises from the informational, with binary distinctions as the deepest layer of substrate from which all structure emerges.

Axiom 1 — The Measurement Imperative

Information must be measured to have meaning. Therefore a measurement substrate must exist.

Axiom 2 — The Sphere as Primitive

Given the Principle of Generative Economy, the unique minimum sufficient measurement substrate is the sphere — the only isotropic primitive requiring zero orientation parameters. The circle is its 2D expression. Hexagonal close-packing and the FCC lattice follow as the unique minimum sufficient tilings of this primitive in 2D and 3D respectively.

Axiom 3 — Signal Conservation

Informational signals are conserved across flops. They cannot be created, destroyed, or deferred.

Axiom 4 — Causal Anchoring

Confirmed states cannot be overwritten. A confirmed state is the geometric record of the signals that produced it — to overwrite it would violate signal conservation. The system therefore has only one available direction: forward. Each flop adds a new layer to the causal record without disturbing prior layers. This directional necessity is the geometric origin of flop time τ .

Summary

All subsequent structures — triadic confirmation, the binary alphabet, the $1 \rightarrow 4$ multiplier, FCC geometry, the conjugate field, the $[[4,2,2]]$ structural correspondence, the identification of the Planck length and Planck time, the derivation of $c_{\text{TMS}} = 1/\sqrt{2}$, and the wave equation — are derived from these axioms and the Principle of Generative Economy without additional assumption.

1.1 Maximum Entropy Stochastic Field (U_{sh})

The ground state of the TMS is U_{sh} as defined in Axiom 0 — an irreducible, algorithmically incompressible binary flux. It is the primordial field of potential from which all structured information is actualized. U_{sh} acts on the informational layer — the latent bits populating the interstices of $M\alpha$. It does not affect the agents, which remain immutable substrate.

1.2 Agent Manifold ($M\alpha$)

A 2D hexagonal array of spherical primitives (α). Agents are the immutable geometric substrate of the TMS. They do not change, move, or degrade. By Axiom 2 and the Principle of Generative Economy, the sphere is the unique primitive — zero orientation parameters, minimal surface-to-volume ratio. Any other primitive introduces orientation degrees of freedom inconsistent with the Principle of Generative Economy.

1.3 Touch Vectors

Each agent α_i maintains six touch vectors connecting it to its six neighbors. Touch vectors are the channels of informational communication between agents. A channel is shared infrastructure — the same vector can carry signals relating to different bits — but each signal is unique to the specific bit it describes. Touch vectors point either laterally within a layer or forward into the next layer. There is no backward propagation. Each signal carried by a touch vector inherits the polarization $\chi \in \{-1, +1\}$ of the confirmation event that produced it.

1.4 Bits, Confirmation, and the Binary Alphabet

Bits (B) exist in the interstices of $M\alpha$ in a latent state, drawn from U_{sh} . By the Principle of Generative Economy, a bit is exactly the size of the triangular interstice formed by three touching spheres — the unique scale at which a bit can be unambiguously localized with complete angular closure. Actualization — the transition of latent B to confirmed $\bar{B}$ — requires triadic agreement among agents.

A confirmed bit carries a binary value:

$$\bar{B} \in \{0, 1\}$$

$\bar{B} = 1$ denotes a confirmed presence — a positive informational result. $\bar{B} = 0$ denotes a confirmed absence — a negative informational result. Both are stabilized outcomes pulled from U_{sh} by triadic confirmation; both lock the interstice into the causal record; both render a confirmed tetrahedral apex at the next layer (Section II). The distinction between the two is not whether the interstice is locked (every confirmed interstice is locked) but what computational content is locked there.

The latent state B in U_{sh} — neither 0 nor 1, but the algorithmically incompressible superposition of both — is ontologically distinct from the confirmed zero $\bar{B} = 0$. A latent bit is unmeasured; a confirmed zero is the measured outcome of an absence. This trinary distinction — latent, confirmed-zero, confirmed-one — is the substrate's irreducible computational state space.

The signed representation:

$$\chi \equiv 2\bar{B} - 1 \in \{-1, +1\}$$

is referred to as the polarization of the confirmed bit and is used in subsequent derivations.

Only bits actualize and deactualize; agents do not. The mechanism by which incoming signals bias the resolution of a confirmation toward $\bar{B} = 1$ or $\bar{B} = 0$ is developed in Volume 1, Paper 2. This paper treats the binary alphabet as an ontological primitive established by the interaction of triadic confirmation with the maximum-entropy substrate.

1.5 The Flop

A flop is the atomic unit of state change in the TMS. It is the complete event sequence initiated by a triadic confirmation: the locking of a bit $\bar{B}$ with binary value, the simultaneous generation of 12 outgoing touch vector signals carrying the polarization of the confirmation, the consumption of those signals as 4 synthetic bits, and the resulting Z-displacement to a new layer. A flop is indivisible — all sub-events occur simultaneously as a single act. There is no sub-flop time. The discrete time parameter τ counts flops. All rates in the TMS — confirmation rate, synthesis rate, propagation speed — are expressions of this single underlying tempo.

These primitives and their interactions define the complete ontology of the TMS. All subsequent structures are derived from them without additional assumption.

II. FCC Necessity

2.1 The Triadic Minimum

Theorem 0.0.2 (Triadic Minimum). *3 agents is the necessary and sufficient minimum for bit actualization in $M\alpha$.*

Proof. The confirmation substrate must constitute a closed simplicial cover of the plane — every interstice fully bounded with complete angular closure. No open edges are permitted; a bit cannot be localized without a closed boundary.

Case 1: 2 agents is insufficient. Two agents define a 1-simplex. A 1-simplex has no interior and subtends less than 360° regardless of arrangement. The enclosure is always open. No bit can be localized. 2-agent confirmation is geometrically impossible.

Case 2: 4 agents is non-minimal. Four agents per interstice produces a closed simplicial cover — the Cartesian square tiling, $\bar{B} = 4\bar{\alpha}_2$. This is geometrically valid but violates the Principle of Generative Economy: it achieves complete angular closure at one degree of complexity greater than necessary.

Case 3: 3 agents is necessary and sufficient. Three agents define a 2-simplex. A 2-simplex achieves complete angular closure — 360° — with no redundancy. Removing any one agent opens the boundary. Adding any agent violates the Principle of Generative Economy. The triangular confirmation geometry tiles the plane completely as a closed simplicial cover, producing the hexagonal lattice as consequence.

Therefore $\bar{B} = 3\bar{\alpha}_2$ is the unique solution selected by the Principle of Generative Economy. The hexagonal lattice is not assumed — it follows from the triadic minimum.

□

$$\bar{B} = 3\bar{\alpha}_2(2D)$$

2.2 Hexagonal Packing

Given the sphere as the unique geometric primitive under Axiom 2, the densest packing of equal discs in 2D is the hexagonal arrangement (Thue, 1890; rigorously proved Fejes Tóth, 1940). This is the unique maximum-density solution and defines the structure of $M\alpha$. It is the direct consequence of the triadic minimum — the hexagonal lattice is what a closed simplicial cover of circles necessarily produces.

2.3 Causal Anchoring and Z-Displacement

When a bit is confirmed in layer S_0 , the confirming agents and their relationships constitute the causal record of that event. By Axiom 4, this record cannot be overwritten. The $1 \rightarrow 4$ multiplier (Section IV) generates synthetic bits that must be resolved at new positions. Three candidate directions exist:

Lateral expansion within $M\alpha$ is geometrically impossible. By Axiom 2 and the Principle of Generative Economy, $M\alpha$ is hexagonally close-packed — the unique maximum-density arrangement of equal spheres in 2D (Thue, 1890). No unoccupied interstice exists at the same layer into which a new agent-scale sphere can be inserted without overlapping an existing agent. Lateral expansion violates Axiom 2.

Oblique displacement introduces a free parameter. Any displacement at an angle other than perpendicular to $M\alpha$ requires specifying that angle — an orientation parameter with no geometric basis in the existing structure. This violates the Principle of Generative Economy.

Orthogonal displacement is geometrically forced. When a new sphere of identical radius settles into the triangular interstice formed by three mutually touching spheres of $M\alpha$, its center comes to rest at a uniquely determined height above the plane — fixed entirely by the contact geometry of equal spheres. This height defines the perpendicular direction. The third dimension is not imported from a pre-existing Euclidean container. It is the unique direction produced by sphere-on-sphere contact in the triangular interstice. Orthogonality is derived, not assumed.

The displacement therefore defines the Axis of Causal Succession:

$$S_0 \rightarrow S_1 : \Delta z = 1$$

Each new layer S_n is the physical record of a confirmed state. This directional necessity is the geometric origin of flop time τ .

2.4 FCC Selection over HCP

When a second layer S_1 forms above S_0 , it occupies the interstices of the first — position set B over position set A. The two candidate stacking sequences for a third layer are ABABAB (hexagonal close-packed, HCP) and ABCABC (face-centered cubic, FCC).

To evaluate these candidates, we must first define what constitutes a distinct position in the causal stack.

Definition 0.0.3 (Causal Position). A position in the TMS causal stack is defined by its complete touch vector connectivity — the full set of geometric relationships between that node and every node in all prior layers. Two positions are causally identical if and only if their connectivity patterns are geometrically indistinguishable to any downstream signal.

This definition is forced by Axiom 4. The causal record is a record of signal history — what arrived, from where, and in what geometric relationship. If two positions share identical connectivity patterns, then a signal reaching either position carries no geometric information about which layer it originated from. The causal record collapses.

Theorem 0.0.4 (FCC Necessity). *Given Definition (Causal Position) and Axiom 4, the unique permitted stacking sequence is ABCABC.*

Proof. In the ABABAB sequence, S_2 returns to A positions. The touch vector connectivity pattern of an A-position node in S_2 is geometrically congruent to that of the corresponding A-position node in S_0 — both sit directly above B-position nodes in an intervening layer, with identical angular relationships to their neighbors. By Definition (Causal Position), these nodes are causally identical. A signal propagating through S_1 into S_2 cannot geometrically distinguish whether it has arrived at a new causal record or returned to S_0 . The causal stack becomes periodic with period 2 — a closed causal loop. This violates Axiom 4: the confirmed state of S_0 is effectively overwritten by the indistinguishability.

In the ABCABC sequence, S_2 occupies C positions — geometrically distinct from both A and B. The touch vector connectivity of every C-position node is unique: it is above a B-position node in S_1 , which is itself above an A-position node in S_0 , producing an asymmetric connectivity pattern that distinguishes S_2 from every prior layer. No two layers share a connectivity pattern. The causal record is monotonically distinguishable at every layer.

HCP is excluded by causal anchoring. FCC is the unique stacking sequence consistent with Axiom 4. □

Corollary 0.0.5 (Sublattice Bifurcation at Recursive Levels). *The FCC structure derived here applies at the substrate level. At each hierarchical level $k \geq 1$ produced by recursive coarse-graining (Paper 3 § 4.5), the meta-FCC sublattice further splits into two interpenetrating FCC sublattices related by spatial inversion through the (1,1,1)-direction offset. The full bifurcated structure that results, and the empirical verification of this bifurcation through $L_{\text{box}} = 12$, is derived in Paper 5 § 6.1.3 and Appendix E. The substrate-level FCC of this paper is the level-0 case of the sublattice-bifurcation recursion that operates at every level $k \geq 1$.*

2.5 Triangular to Tetrahedral Interstices

The ABCABC stacking sequence converts triangular interstices bounded by 3 agents into tetrahedral interstices bounded by 4 agents. The confirmation geometry is preserved across the dimensional transition:

$$\bar{B} = 3\bar{\alpha}_2(2D)$$

$$\bar{B} = 4\bar{\alpha}_3(3D)$$

The triadic rule does not change. The geometry scales naturally from triangle to tetrahedron without introducing any new mechanism. The binary alphabet $\bar{B} \in \{0, 1\}$ of Section 1.4 is preserved unchanged: both 2D and 3D confirmation events resolve to a binary value.

The tetrahedral confirmation unit — four agents enclosing one bit with one of two possible values — emerges as the fundamental 3D structure. Its informational dynamics are derived in Section IV.

III. The Fundamental Length Scale

With the FCC geometry established, the system possesses a unique length scale — the agent diameter. This section shows that scale is not a free parameter but is forced to equal the Planck length, and derives the Planck time as a necessary consequence.

3.1 Definition

Definition 0.0.6 (Fundamental length). Let ℓ_0 denote the diameter of the spherical primitive α . This is the unique length scale appearing in the definition of $M\alpha$.

3.2 Planck Length Identification

Theorem 0.0.7 (Planck Length Identification). *The fundamental agent spacing ℓ_0 is identical to the Planck length ℓ_p . This identification carries no free parameters.*

Proof. The 2D manifold $M\alpha$ is the complete pre-geometric substrate — no structure exists below it. Any measurement of distance in the emergent 3D geometry must ultimately resolve to positions of agents or interstices within $M\alpha$. Two distinct agents cannot occupy positions closer than ℓ_0 , as their spheres would overlap, violating Axiom 2. Therefore ℓ_0 is the minimum resolvable length in the emergent 3D geometry. Below ℓ_0 , no geometric distinctions exist — only the 2D substrate.

The Planck length $\ell_p = \sqrt{\hbar G/c^3}$ is the scale below which classical spacetime geometry ceases to be meaningful. The TMS provides the geometric reason: below ℓ_0 , the emergent 3D geometry does not exist. The identification

$$\ell_0 \equiv \ell_p$$

is therefore forced by the ontology. It is not an external input — it is the statement that the pre-geometric substrate sets the boundary of geometry. No free parameter remains. □

Remark 0.0.8 (On the Status of $\hbar$ and G). The identification $\ell_0 \equiv \ell_p$ fixes the substrate's minimum length, and § 3.3 fixes its tick τ_{flop} ; from these the speed c is derived. It is worth stating precisely what this leaves, because the status of $\hbar$ and G is a point of principle, not a deferred calculation.

The substrate is combinatorial: it produces dimensionless quantities — counts, ratios, angles (the coordination number 12, the subshell capacities 2, 6, 10, 14, the noble-gas atomic numbers, the fermionic sign, the reconstructed angles of Paper 5 § 4.1). A dimensionful constant cannot be produced by dimensionless structure alone; it requires a unit, which is scale-setting data, not a number the combinatorics can output. The three constants $\{c, \hbar, G\}$ form an invertible basis for the three mechanical dimensions (mass, length, time) — fixing them is exactly the choice of Planck units — so any theory expressed in dimensionful units must be handed a scale for each dimension it cannot supply intrinsically.

The TMS accounting is favorable and exact. The substrate supplies *two* of the three scales from its own structure: the length ℓ_p (lattice geometry) and the time τ_{flop} (the update tick). Their ratio gives the third mechanical quantity of the length–time sector, the speed c (§ 3.3), so c is derived rather than imported. What remains is the mass/action sector. The substrate's fundamental quantum is a *bit* — dimensionless information — not an action; a combinatorial substrate therefore has no intrinsic action scale, and exactly one dimensionful

input must be supplied to reach physical units. That single input is what $\hbar$ carries. Newton's constant is then not independent: given ℓ_p , c , and $\hbar$, the Planck relation $\ell_p = \sqrt{\hbar G/c^3}$ inverts to $G = c^3 \ell_p^2 / \hbar$, so G is *derived*. The framework thus takes exactly one dimensionful scale-setting input, $\hbar$, and yields c and G — one import, not two open constants, and the minimum any theory stated in dimensionful units can require.

This sharpens, rather than weakens, the no-free-parameters claim. TMS has no free *dimensionless* parameters: every dimensionless quantity is fixed by the axioms. It has exactly one dimensionful input, $\hbar$, which sets the scale at which the dimensionless structure is read — and even that does not float freely, since it fixes c and G with it. Asking the framework to output the numerical *value* of $\hbar$ in SI would be a dimensional category error, equivalent to asking it to output the metre or the second: that value is the definition of the unit system, not a fact the combinatorics decides. This is the same footing on which the Standard Model fixes its group-theoretic structure from first principles while taking its dimensionful couplings from experiment.

3.3 Flop Time

Corollary 0.0.9 (Flop Time). *The substrate has two intrinsic dimensionful scales — the minimum length ℓ_p (§ 3.2) and the flop time τ_{flop} , the duration of a single update. The propagation speed is the ratio of the two, and the flop time is identified with the Planck time:*

$$c = \frac{1}{\sqrt{2}} \frac{\ell_p}{\tau_{\text{flop}}}, \quad \tau_{\text{flop}} = t_p$$

Proof. The flop is the substrate's irreducible tick: there is no sub-flop time (§ II), and the discrete time parameter τ counts flops. The flop time τ_{flop} is therefore an intrinsic primitive of the dynamics, not a quantity constructed from other scales. Together with the intrinsic length ℓ_p , it fixes a speed: a signal advances one face-diagonal step per flop, giving $c = (1/\sqrt{2}) \ell_p / \tau_{\text{flop}}$ in lattice units (the $1/\sqrt{2}$ is the FCC face-diagonal factor derived in § VI). The speed of light is thus an *output* of the substrate — the lattice's own steps-per-tick rate — not an external input. The wave equation of § VI confirms the identification: the physical wave speed from the discrete dynamics equals c precisely when the flop time is the Planck time, $\tau_{\text{flop}} = t_p = \sqrt{\hbar G/c^5}$. One flop equals exactly one Planck time. $\square$

3.4 Agent Density

Corollary 0.0.10 (Agent Density). *The agent density in $M\alpha$ is $2/(\sqrt{3} \cdot \ell_p^2)$ per unit area. The FCC lattice inherits the same fundamental spacing ℓ_p between agent centers.*

IV. The Triadic Actualization Rule and the $1 \rightarrow 4$ Multiplier

4.1 The Triadic Constraint

The minimum boundary condition for bit actualization in $M\alpha$ is a 2-simplex (triangle) formed by three agents, as derived in Section II. Each agent has six touch vectors. Each bit costs exactly three unique signals — one from each vertex of the confirming triangle. An agent may participate in up to six confirmations (one per neighbor pair), but each confirmation consumes three unique signals. No signal describes more than one bit.

$$\bar{B} = 3\bar{\alpha}_2$$

4.2 The Multiplier Proof

Given a primary confirmation event S_0 involving agents $\{\bar{\alpha}_1, \bar{\alpha}_2, \bar{\alpha}_3\}$:

- Each confirming agent has 6 touch vectors
- 2 vectors per agent are locked into the confirmation triangle
- 4 vectors per agent carry the signal outward
- Total outgoing signals: $3 \times 4 = 12$

Since each bit costs exactly three unique signals:

$$12 \div 3 = 4 \text{ synthetic bits}$$

This division is exact by necessity. The hexagonal geometry is the unique 2D packing where 6 neighbors per agent produces a clean integer result under the triadic constraint. No remainder is geometrically permitted — any configuration that would produce one is not a valid confirmation geometry.

4.3 Informational Closure and Polarization Transport

Every confirmation event is informationally complete within a single flop. All 12 outgoing signals are fully consumed by the 4 synthetic bits. Nothing is stored, deferred, or lost. The TMS is informationally closed by construction — a direct consequence of Axiom 3. Closure here is of signal count and content: every signal is consumed, none created or destroyed (Axiom 3). What the signed aggregate does not preserve is resolvable distinction — three incoming polarizations collapse to one signed outcome — and this per-flop loss of distinction, not of content, is the geometric origin of the irreversibility developed in Paper 5.

The 12 outgoing signals carry the polarization of their source confirmation. A $\bar{B} = 1$ confirmation emits 12 signals carrying $\chi = +1$; a $\bar{B} = 0$ confirmation emits 12 signals carrying $\chi = -1$. The signal count is identical in both cases — informational closure operates on count alone — but the polarization content differs.

The $1 \rightarrow 4$ multiplier is content-preserving: each downstream synthetic bit position accumulates the polarizations of its 3 incoming signals as a signed aggregate. The resulting polarization landscape across the synthetic bit positions of S_{n+1} biases the subsequent resolution of those positions toward $\bar{B} = 1$ or $\bar{B} = 0$ against the unbiased stochastic kick from U_{sh} . The propagation dynamics of this polarization landscape, the per-flop biasing mechanism, and the structural correspondence to the Born rule are developed in Volume 1, Paper 2.

The informational closure of each confirmation event, combined with the $1 \rightarrow 4$ multiplier, generates a structural surplus that cannot be resolved within the existing 2D layer — necessitating the Z-displacement and FCC geometry derived in Section II. The stability properties of the resulting tetrahedral structure are examined in Section V.

V. The $[[4,2,2]]$ Structural Correspondence and Fault-Tolerant Propagation

5.1 The Stabilizer Mapping

The $[[4,2,2]]$ quantum error-detection code encodes 2 logical qubits in 4 physical qubits with code distance 2. Its stabilizer group is generated by:

$$S_x = X_1 X_2 X_3 X_4$$

$$S_Z = Z_1 Z_2 Z_3 Z_4$$

A state lies in the code space when both stabilizers evaluate +1. Any single-qubit error anticommutes with at least one stabilizer, producing a measurable syndrome (the stabilizer evaluates to -1). The two-dimensional code space encodes 2 independent logical qubits — 4 distinguishable logical states.

The tetrahedral confirmation unit of the TMS maps onto this structure as follows:

- 4 physical qubits $\rightarrow$ 4 agents $\{\bar{\alpha}_1, \bar{\alpha}_2, \bar{\alpha}_3, \bar{\alpha}_4\}$ of the tetrahedral interstice.
- Logical qubit 1 $\rightarrow$ the value of the confirmed bit in the current plane: $\bar{B} \in \{0, 1\}$.
- Logical qubit 2 $\rightarrow$ the inheritance of that value through the causal stack — the integrity of the bit's history preserved across prior layers.

Every confirmed bit therefore carries two independent informational degrees of freedom: its current value and its causal history. The 4 logical states of the $[[4,2,2]]$ code space correspond exactly to the 4 possible combinations of these two qubits. The encoding is faithful.

Single-agent disruptions caused by U_{sh} fall into two algebraically distinct types, each detected by a different stabilizer:

- Bit-flip-type disruption (X-type error). A disturbance that would alter the value contribution of a single agent. This anticommutes with the Z-stabilizer, producing $S_Z = -1$.
- Phase-type disruption (Z-type error). A disturbance that would corrupt the inheritance contribution of a single agent. This anticommutes with the X-stabilizer, producing $S_X = -1$.

Both syndrome flavors are checked at every flop. The confirmation mechanism that pulls a bit from U_{sh} is the triadic actualization rule of Section II; the $[[4,2,2]]$ stabilizers do not perform confirmation but rather verify the integrity of already-confirmed bits against subsequent disruption. The dynamic-code framing of recent work on stabilizer codes emerging from spacetime circuits (Aitchison & Béri, 2025) provides a parallel formal context for the structural correspondence developed here.

5.2 The Restoring Force

When U_{sh} disrupts a confirmed bit $\bar{B}$ — flipping its value or corrupting its causal history — the 4 agents of the tetrahedral unit are unchanged. They remain in place as immutable substrate. The confirmation relationship between them is violated, producing a detectable parity syndrome.

The synthetic bits produced at the current flop — the same flop at which the violation is detected — repair the missing or corrupted confirmation, restoring parity. Detection and repair are a single atomic act, a direct consequence of informational closure. The system cannot defer repair any more than it can defer signal consumption. Everything resolves within the flop.

Both logical qubits are detected and repaired within the same flop. A bit-flip-type violation ($S_Z = -1$) and a phase-type violation ($S_X = -1$) are independent failure modes, protected independently. In every flop, the TMS simultaneously maintains the present state, the binary value of every confirmed bit, and the integrity of the entire causal record. The system is not merely propagating forward — it is actively preserving the full informational content of its history as a single geometric motion.

5.3 Code Distance and Fault-Tolerant Propagation

The $[[4,2,2]]$ code has distance 2 — any single error is detectable. In TMS terms:

- A single bit-flip-type disruption on one agent violates S_Z while leaving S_X intact — detectable.
- A single phase-type disruption on one agent violates S_X while leaving S_Z intact — detectable.
- The tetrahedral geometry guarantees that no single-agent disruption can leave both stabilizers intact — code distance 2 follows from the geometry alone.

The full stabilizer algebra — both stabilizers, both logical qubits, the assignment of value and history to the two logical degrees of freedom, and the code distance — is derived from TMS primitives. The mapping constitutes a structural correspondence. The framework is also consistent with the broader holographic quantum error correction tradition (Pastawski, Yoshida, Harlow, & Preskill, 2015; El Morsalani, 2025) that has informed the use of stabilizer codes in emergent-spacetime contexts. $\square$

The $[[4,2,2]]$ structural correspondence guarantees that the TMS substrate is fault-tolerant against stochastic interference from U_{sh} . Confirmed bits propagate across the lattice with their values and histories intact, because the error-detection property of each tetrahedral unit preserves both at every hop. This establishes the TMS lattice as a stable propagation medium for both the confirmation field and the polarization landscape carrying binary content.

The discrete dynamics of this fault-tolerant lattice, when taken to the continuum limit, reproduce the standard wave equation — derived in Section VI.

The structural correspondence developed here applies at the substrate level. The same $[[4,2,2]]$ stabilizer structure reappears at every hierarchical level by axiom scale-independence (Paper 3 § 4.7), with one structural addition at level $k \geq 1$: the meta-FCC splits into two interpenetrating sublattices A and B (Paper 3 § 4.5; Paper 5 § 6.1.3), each running independent $[[4,2,2]]$ dynamics. The level-1 verification is empirical (Paper 5 Appendix E, Stages 5 and 6). The cross-level commutator structure between the A-sublattice and B-sublattice stabilizers at level k provides the substrate-level origin of non-Abelian gauge closure at level $k + 1$ (Paper 5 § 8.2).

5.4 Scale-Invariant Stabilizer Null-Space

The $[[4,2,2]]$ stabilizer structure of § 5.1–§ 5.3, applied to an FCC patch of finite extent L_{box} at any hierarchical level, has a precise count of stabilizer-satisfying configurations: the dimension of the null space over F_2 of the tetrahedron-to-site parity matrix. We state this as a theorem.

Theorem 0.0.11 (Scale-Invariant Stabilizer Null-Space). *For an FCC patch of edge L_{box} at any hierarchical level, with $[[4,2,2]]$ stabilizers imposed at every tetrahedron in the patch, the dimension of the stabilizer-satisfying configuration space over F_2 is $k = 3n$, where $n = (L_{\text{box}} - 1)/2$ at level 1 and analogously at higher levels (the level-local lattice depth). The number of distinct stabilizer-satisfying configurations is therefore $2^{(3n)}$, and this scaling law is a structural property of the FCC + $[[4,2,2]]$ dynamics that is independent of hierarchical level.*

Proof. The $[[4,2,2]]$ stabilizers at each tetrahedron act linearly on the F_2 – valued polarization state of the four agents in the tetrahedron. Their joint action on a patch of N sites with T tetrahedra is a linear map $F_2^N \rightarrow F_2^T$ whose kernel dimension is the satisfying configuration count. The geometric argument that this kernel dimension equals $3n$ at every level is given in schematic form by the FCC cubic-volume scaling: the unit cubic cell of side 2 in level-local lattice units contributes 3 degrees of freedom to the null space (one per

axis-pair of the cubic cell), and a patch of cubic-cell depth n contains n^3 unit cells along with their bounding stabilizer relations, giving $3n$ dimensions of freedom from the cubic-cell-edge degrees and $(n^3 - 3n)$ parity constraints from the tetrahedron-interior stabilizers. The result is $k = 3n$. The empirical verification of this law at substrate level ($n = 1$ through L_{box} extents matching the patch sizes) and at levels 1 and 2 (Paper 5 Appendix E, Stages 5 and 6) confirms $k = 3n$ through $n = 6$ across all levels tested.

□

Remark 0.0.12. The scale-invariance of the $k = 3n$ law is the structural reason the recursion-preservation theorem of Paper 5 § 2.5 has its full force: the $[[4,2,2]]$ dynamics at every level have the same combinatorial structure when measured in level-local units, because the null-space scaling depends only on the FCC patch geometry, not on the absolute scale.

VI. The Wave Equation

6.1 The Confirmation Field

Let $\varphi(\bar{r}, \tau)$ denote the local confirmation density — the scalar field representing the density of actualized bits $\bar{B}$ at any point in the lattice:

$$\varphi(\bar{r}, \tau) = \text{confirmedbitsinregion} / \text{totalavailableintersticesinregion}$$

This field counts confirmation events without regard to their binary value. It is purely informational and purely geometric. It requires no additional physical assumptions.

6.2 The Discrete Laplacian on the FCC Lattice

Each node in the FCC lattice has 12 nearest neighbors. The discrete Laplacian is defined as:

$$\Delta_{\text{FCC}} \varphi_i = \Sigma(\varphi_j - \varphi_i), j \in \text{neigh}(i)$$

This operator measures the difference between the confirmation density at node i and its neighbors — the local informational gradient across the lattice.

6.3 The Conjugate Field and Second-Order Dynamics

Lemma 0.0.13 (Field Conjugacy). *The confirmation field φ and the forward-propagating excitation field Π constitute a co-generated conjugate pair, arising necessarily from the interaction between the $1 \rightarrow 4$ cascade and U_{sh} regulation.*

Proof. Step 1 — The 2D case produces first-order closure. In the 2D manifold $M\alpha$ without Z-displacement, all signals generated by a confirmation event are consumed within the same layer. The state at any layer is fully described by φ alone. No forward-propagating field persists. The update rule is first-order by construction — confirmation density at the next step depends only on confirmation density at the current step.

Step 2 — The $1 \rightarrow 4$ cascade generates exponential confirmation pressure. In the 3D FCC lattice, each confirmation event at layer S_n generates 12 outgoing signals, consumed as 4 synthetic confirmations at S_{n+1} . Each of those 4 confirmations generates 12 signals, consumed as 16 confirmations at S_{n+2} . Left unchecked, this produces an exponential cascade:

$$1 \rightarrow 4 \rightarrow 16 \rightarrow 64 \rightarrow \dots$$

This cascade is not a pathology — it is the direct geometric consequence of informational closure applied layer-to-layer under FCC connectivity. Every signal is consumed. Every consumption generates new signals. The system is structurally amplifying.

Step 3 — U_{sh} regulation produces dynamic stability. U_{sh} acts on the latent bit population at every layer, disrupting confirmations stochastically and returning bits to the unactualized field. This regulation is the only mechanism that competes with the $1 \rightarrow 4$ cascade.

The cascade generically outpaces disruption because amplification is deterministic and geometric while disruption is random. At each layer, the $1 \rightarrow 4$ mechanism advances with fixed multiplicity; U_{sh} disrupts with probability distributed across the full bit population. On average, the cascade stays ahead — not because the system reaches equilibrium, but because geometric amplification is structurally faster than uncoordinated stochastic interference.

The $[[4,2,2]]$ fault tolerance (Section V) reinforces this at the local level. Each tetrahedral confirmation unit detects and repairs single disruptions within the same flop, preserving causal coherence node by node. As the structure grows, its surface area exposed to U_{sh} increases — but so does the interior volume protected by the causal stack. The propagating wave is the coherent interior structure advancing faster than disruption can reach it.

This is dynamic stability, not static equilibrium. The system does not settle — it persists by continuing to move forward, in the same way a bicycle is stable under motion but not at rest. A sufficiently large simultaneous disruption — a true stochastic black swan event collapsing enough bits at once to exceed the current flop's resolution capacity — would constitute a system-wide failure. Such a failure produces no downstream causal layer and therefore leaves no record within the stack. By Axiom 4, only what is confirmed is recorded. A collapsed system is unobservable from within by construction.

The propagating residual Π is therefore the net forward-propagating confirmation pressure in a dynamically stable instance — the touch vector field carrying well-defined inherited excitation forward across layers under the constraint of Axiom 3.

Step 4 — Signal conservation forces the coupling equations. At any layer S_n , the system state requires two quantities for complete specification:

- $\varphi(S_n)$ — the confirmation density: what actually locked at this layer
- $\Pi(S_n)$ — the inherited excitation pressure: what is being driven forward into S_{n+1}

By Axiom 3, signals are conserved across flops. The rate of change of confirmation density is exactly the inherited excitation pressure — no signal is created or destroyed in the transfer:

$$\partial\varphi/\partial\tau = \Pi$$

The rate of change of excitation pressure is governed by the local confirmation gradient across the FCC lattice:

$$\partial\Pi/\partial\tau = c^2\Delta_{\text{FCC}}\varphi$$

These equations are not assumed. They are the unique expressions of Axiom 3 applied to a two-quantity state under FCC connectivity.

□

Remark 0.0.14 (Uniqueness of the Coupling). The second coupling equation $\partial\Pi/\partial\tau = c^2\Delta_{\text{FCC}}\varphi$ admits no additional local term in φ under the FCC touch-vector geometry. A candidate alternative of the form $\partial\Pi/\partial\tau = c^2\Delta_{\text{FCC}}\varphi + \alpha\varphi$ would require that the forward-propagating excitation pressure at a node be modified by a quantity proportional to the confirmation density *at that same node*, with no reference to its neighbors. This is a self-amplification term: the node contributes to its own forward pressure without communicating through the touch vector network.

Such a term is forbidden by the substrate’s connectivity architecture. By Paper 1 § 1.3, the touch vectors are the *only* channels of informational communication in the TMS — every signal that updates a quantity at a node arrives across a touch vector from a neighbor. There is no on-site signaling channel: an agent does not communicate with itself, and a bit cannot bias its own forward propagation independently of its neighbors. The discrete Laplacian Δ_{FCC} is the unique nearest-neighbor local operator consistent with this constraint — it sums differences $(\varphi_j - \varphi_i)$ across the 12 face-diagonal neighbors and is identically zero for spatially uniform φ . An $\alpha\varphi$ term, by contrast, is non-zero for uniform φ , which would imply that a spatially uniform confirmation density generates forward propagation in the absence of any spatial gradient — propagation from no source.

Axiom 3 (signal conservation) closes this. Signals propagate; they do not spontaneously generate. A forward pressure that responds to local density rather than local gradients would represent a creation of propagation content with no corresponding signal expenditure across the touch vector network. The $\alpha = 0$ case is therefore the unique selection consistent with both nearest-neighbor connectivity and signal conservation.

The same argument forbids non-local operators (which would require channels beyond nearest-neighbor touch vectors), fractional-derivative operators (which would require memory of confirmation density across multiple prior layers, violating the atomicity of the flop), and higher-derivative spatial operators beyond Δ_{FCC} (which would require coupling across more than nearest-neighbor distance per flop). The discrete Laplacian Δ_{FCC} acting on φ alone is the unique local operator consistent with the touch vector architecture; its coefficient c^2 is fixed by the propagation-speed derivation of § 6.5. The second coupling equation is therefore forced in both functional form and coefficient. $\square$

Step 5 — Elimination yields second-order dynamics. Eliminating Π from the coupled system:

$$\partial^2\varphi/\partial\tau^2 = c^2\Delta_{\text{FCC}}\varphi$$

The second-order structure is not imported from classical field theory. It is the necessary consequence of a self-amplifying confirmation cascade regulated by a maximum-entropy stochastic field. A 2D confirmation has no cascade, no inherited pressure, and no inertia. The transition to 3D, forced by the $1 \rightarrow 4$ multiplier, introduces exactly the two-quantity state that produces wave rather than diffusive behavior. $\square$

This result is structurally analogous to the Hamiltonian formulation of classical field theory, where conjugate pairs (q, p) produce second-order equations of motion via elimination of momentum (Goldstein, Poole & Safko, 2002). The TMS recovers this structure from the population dynamics of a self-amplifying confirmation cascade under stochastic regulation — grounding the conjugate pair (φ, Π) in the geometric and informational architecture of the substrate itself.

6.4 The Discrete Wave Equation

Each confirmation event propagates through the lattice via the touch vector network. The rate of change of φ_i is governed by the local gradient and the fault-tolerant restoring force established in Section V:

$$\partial^2 \varphi_i / \partial \tau^2 = c_{\text{TMS}}^2 \Delta_{\text{FCC}} \varphi_i$$

where τ is the discrete flop time and c_{TMS} is the propagation speed in lattice spacings per flop.

6.5 The Propagation Speed

Theorem 0.0.15 (Propagation Speed). *The unique propagation speed consistent with FCC geometry, the Planck length identification of Section III, and the physical wave speed c is $c_{\text{TMS}} = 1/\sqrt{2}$ lattice spacings per flop.*

Proof. In the FCC lattice with nearest-neighbor distance ℓ_p , the 12 nearest neighbors of any node lie at vectors of the form $(\pm \ell_p/\sqrt{2}, \pm \ell_p/\sqrt{2}, 0)$ and cyclic permutations — all directed along face diagonals, each of magnitude ℓ_p .

Taylor-expanding φ over all 12 neighbors and summing, the first-order terms cancel by the symmetry of the FCC lattice. Computing the second-order tensor over the 12 neighbor vectors r_j :

$$\sum_j r_j \mu r_j \nu = 4\ell_p^2 \delta \mu \nu$$

The discrete Laplacian therefore converges as:

$$\Delta_{\text{FCC}} \varphi \rightarrow 2\ell_p^2 \nabla^2 \varphi$$

Substituting into the discrete wave equation and converting to physical time $t = \tau \cdot \tau_{\text{flop}}$:

$$\partial^2 \varphi / \partial t^2 = (c_{\text{TMS}}^2 \cdot 2\ell_p^2 / \tau_{\text{flop}}^2) \nabla^2 \varphi$$

Requiring this to equal $c^2 \nabla^2 \varphi$, and substituting $\tau_{\text{flop}} = \ell_p / c$ from Section 3.3:

$$c_{\text{TMS}} = 1/\sqrt{2}$$

The factor $1/\sqrt{2} = \cos(45^\circ)$ is the direct geometric projection factor of FCC face-diagonal connectivity. The propagation speed is not a parameter — it is the cosine of the lattice's canonical angle. No free parameter remains.

□

Corollary 0.0.16 (Consistency). *Substituting $c_{\text{TMS}} = 1/\sqrt{2}$ into the physical wave speed relation confirms:*

$$c_{\text{physical}} = c_{\text{TMS}} \cdot \sqrt{2} \cdot \ell_p / \tau_{\text{flop}} = (1/\sqrt{2}) \cdot \sqrt{2} \cdot c = c$$

6.6 The Continuum Limit

Substituting $c_{\text{TMS}} = 1/\sqrt{2}$ and the Laplacian result into the discrete wave equation:

$$\partial^2 \varphi / \partial \tau^2 = (1/2) \cdot 2\ell_p^2 \nabla^2 \varphi = \ell_p^2 \nabla^2 \varphi$$

Converting to physical time $t = \tau \cdot t_p$ and using $\ell_p = c \cdot t_p$:

$$\partial^2 \phi / \partial t^2 = c^2 \nabla^2 \phi$$

Remark 0.0.17 (Stochasticity and the Continuum Limit). U_{sh} operates at every flop as both the source of latent bits available for confirmation and the stochastic disruptor of previously confirmed states. These dual roles are not in tension — they are symmetric expressions of the same maximum entropy property established in Axiom 0. U_{sh} has no attractors and no preferred targets. Every bit in every region has equal probability of disruption at every flop. No geometric bias exists in the disruption mechanism itself.

At the substrate level, U_{sh} contributes two statistically symmetric pressures to the confirmation field: a sourcing pressure drawing new bits toward actualization, and a disrupting pressure returning confirmed bits to the unactualized field. Because both pressures are drawn from the same maximum entropy distribution with no preferred scale or direction, their net contribution to the large-scale confirmation field averages to zero across sufficiently many flops.

What survives to the continuum level is therefore not a noisy wave equation with a stochastic forcing term — it is the structured confirmation cascade, which is the only process with geometric coherence across scales. The wave equation describes this coherent structure. The stochastic substrate does not add a forcing term to the continuum limit because its two roles cancel symmetrically. The continuum field ϕ is the large-scale average of a dynamically stable cascade instance. This observable universe represents one such stable average over flops — not the unique or guaranteed outcome, but the instance whose causal stack persists long enough to be inhabited.

6.7 The Wave Equation

Recognizing flop time τ as the emergent time coordinate t :

$$\partial^2 \phi / \partial t^2 - c^2 \nabla^2 \phi = 0$$

This is the standard wave equation. It is not imposed on the TMS — it emerges necessarily from the discrete dynamics of a fault-tolerant FCC confirmation lattice in the continuum limit. Physical wave phenomena are the large-scale behavior of the confirmation field propagating through the agent substrate.

6.8 The Polarization Field as Parallel Object

Remark 0.0.18 (The Polarization Field). The confirmation density ϕ derived above is one of two independent informational fields supported by the FCC lattice. The polarization field $\psi \equiv \langle \chi \rangle$ — the local coarse-grained mean of confirmed bit values, ranging over $[-1, +1]$ — propagates by an analogous discrete wave equation on the same FCC lattice:

$$\partial^2 \psi / \partial \tau^2 = c_{TMS}^2 \Delta_{FCC} \psi$$

The propagation speed $c_{TMS} = 1/\sqrt{2}$ is inherited from the lattice geometry; the proof of Section 6.5 applies unchanged because both fields propagate across the same nearest-neighbor structure. Both fields share the same fault-tolerant FCC substrate. The two fields evolve independently in the bulk and couple at confirmation events, where the local incoming polarization wave amplitude biases the resolution of new confirmations toward $\bar{B} = 1$ or $\bar{B} = 0$ against the unbiased stochastic kick from U_{sh} .

The dynamics of ψ , its coupling to ϕ at confirmation events, the emergence of computation from polarization-modulated propagation, the structural correspondence between this coupling and the Born rule, and the formal selection criterion for stable subsystem programs are developed in Volume 1, Paper 2.

VII. Predictions and Directions for Subsequent Work

7.1 The Generalized Uncertainty Principle: A Geometric Grounding

The Generalized Uncertainty Principle (GUP) extends the standard Heisenberg relation by introducing a minimum measurable length:

$$\Delta x \cdot \Delta p \geq \hbar/2 \cdot (1 + \beta(\Delta p)^2/m_p^2 c^2)$$

where β is a dimensionless deformation parameter characterizing the scale at which quantum gravitational effects become significant. The GUP has been independently motivated by string theory, loop quantum gravity, and quantum geometry (Hossenfelder, 2013). In all existing formulations, β is a free parameter — inserted phenomenologically rather than derived from geometric first principles.

The TMS provides a geometric grounding for the minimum length that the GUP encodes. The identification $\ell_0 \equiv \ell_p$ established in Section III has a precise geometric meaning: below the agent scale, the emergent 3D geometry does not exist. This is not a statement about measurement difficulty or quantum fluctuations — it is a statement about substrate architecture. The floor on position resolution is not a statistical lower bound but a geometric boundary. At ℓ_p , the continuous confirmation field ϕ approaches its binary alphabet limit — $\bar{B} \in \{0, 1\}$ — the irreducible primitive of the TMS ontology. Spatial fuzziness at the Planck scale is not noise added to geometry. It is the geometry of the substrate itself becoming directly visible.

Examining a bit at the Planck scale means encountering not merely its current confirmation state, but the full geometric and polarization record of how it arrived there — the layered causal history encoded in the down-vector structure. What standard quantum mechanics describes as wavefunction complexity is, in TMS terms, the geometric depth and polarization complexity of the confirmation record. The uncertainty is not epistemic. It is the irreducible causal narrative of the bit made visible at the substrate boundary.

This geometric grounding is consistent with recent independent work deriving GUP-type relations from spacetime microgeometry (Giné, 2026), which proposes that generalized uncertainty principles may emerge as effective kinematical relations induced by microscopic spacetime geometry. The TMS provides an independent derivation from a more fundamental starting point — geometric primitives prior to spacetime, rather than horizon microstructure within it.

7.2 A Novel Prediction: Causal-Depth-Dependent Uncertainty

The TMS architecture implies a specific and novel extension of the standard GUP that, to our knowledge, has no analog in existing formulations.

In the TMS, every confirmed bit exists at the terminus of a causal chain encoded in the down-vector structure. The confirmation field ϕ at any point in the lattice carries not only the current confirmation state but the full geometric record of how that state was reached — the layered history of prior confirmation events that produced the current position in the causal stack. With the binary alphabet established in Section 1.4, each layer of that history additionally carries a specific polarization value. We term the depth of this record the causal depth d of the bit, and the specific pattern of polarizations across that depth the polarization complexity of its history.

We conjecture that the effective deformation parameter β is not a universal constant but a function of both:

$$\beta = f(d, K_\psi)$$

where d is the causal depth and K_ψ is the algorithmic complexity of the polarization landscape encoding the bit's confirmation history. f is monotonically increasing in both arguments. The precise form of f is a prediction of the FCC confirmation geometry and is left for derivation in subsequent work.

The physical interpretation is direct. A recently confirmed bit — shallow causal history, few prior layers, simple polarization record — has a substrate footprint close to the bare agent scale ℓ_p . Its position uncertainty is close to the geometric minimum. A bit embedded in a deep, richly polarized causal structure has a larger substrate footprint — more of the lattice geometry and binary content is geometrically present in its confirmation record. Its effective position uncertainty is correspondingly larger than the bare minimum.

In standard quantum mechanical language: wavefunction complexity scales with both causal depth and polarization complexity, and minimum position uncertainty scales with wavefunction complexity. The TMS provides a geometric derivation of this relationship from the down-vector structure of the FCC causal stack and the binary alphabet recorded within it.

This prediction is testable in principle. If β scales with the joint causal-depth and polarization-complexity profile, then quantum systems with richer interaction histories — more degrees of freedom, deeper entanglement structure, and more complex binary records — should exhibit measurably larger GUP corrections than simpler systems at equivalent energy scales. Recent work on GUP phenomenology in macroscopic quantum systems and gravitational wave detectors (Hossenfelder, 2013) provides experimental frameworks within which causal-depth and polarization-complexity dependence could in principle be distinguished from a universal β .

7.2a The Computational Irreducibility Connection

The causal-depth dependence of β admits a deeper interpretation through the lens of computational irreducibility (Wolfram, 2002). A computationally irreducible system is one for which no shortcut exists to determine its future state — the only way to know where the system is is to trace every step of its history.

In TMS terms, a bit with deep causal history and complex polarization record is computationally irreducible in precisely this sense. The down-vector chain encoding its confirmation record cannot be compressed — there is no shorter description of its current state than the full sequence of prior confirmation events, including the specific binary values resolved at each layer. This is not a limitation of our knowledge. It is a property of the causal record itself.

The connection to Axiom 0 is direct. U_{sh} is defined as algorithmically incompressible with Kolmogorov complexity $K(n) \approx n$ — the maximum entropy ground state from which all structure is actualized. Confirmation events produce locally reducible structure by carving order from that incompressibility. But the record of how that order was produced — the down-vector chain with its specific polarization sequence — inherits the incompressibility of its source. Deep causal stacks with complex polarization landscapes are not merely geometrically large. They are computationally irreducible in the strict Kolmogorov sense.

The effective position uncertainty $\Delta x \geq \ell_p \cdot f(d, K_\psi)$ therefore has a computational interpretation: measuring the position of a deeply embedded bit to better than $f(d, K_\psi) \cdot \ell_p$ would require compressing an irreducible computation. This is impossible by definition — not in principle, but as a consequence of the same Kolmogorov incompressibility that characterizes U_{sh} itself. The GUP minimum length in TMS is not

merely a geometric floor. It is the Planck-scale expression of computational irreducibility applied to position measurement.

This connects the paper's opening axiom — the incompressibility of U_{sh} — directly to its closing prediction. The arc from Axiom 0 to the GUP is a single unbroken chain:

$$U_{\text{sh}} [K(n) \approx n] \rightarrow \text{confirmation events} \rightarrow \text{binary alphabet} \rightarrow \text{polarization-rich causal stack} \rightarrow \beta = f(d, K_{\psi}) \rightarrow \text{GUP}$$

7.3 A Unified Geometric Account of Decoherence

The same causal depth structure that predicts β — scaling offers a geometric account of quantum decoherence. In standard quantum mechanics, decoherence time depends on the number of environmental degrees of freedom — macroscopic objects decohere almost instantaneously because they interact with enormously many environmental modes. The mechanism is understood phenomenologically but lacks a geometric derivation from first principles.

In the TMS, the causal stack provides a natural geometric analog. A bit with shallow causal depth has a small substrate footprint — it interacts with few prior confirmation layers, maintains a simple touch vector structure, and is relatively isolated from the full lattice history. Such a bit maintains quantum coherence longer because its geometric entanglement with the causal record is limited.

A bit with deep causal history has an extensive substrate footprint — its confirmation record spans many layers, its touch vector structure is richly connected to the broader lattice, and it is geometrically embedded in the full causal record. Such a bit decoheres rapidly because its geometric connection to the lattice history is extensive — there are many more channels through which U_{sh} disruption can reach and collapse its coherent state.

Decoherence time, in this picture, is inversely proportional to causal depth. Macroscopic objects decohere almost instantaneously because they are geometrically ancient — their causal stack is deep beyond measure. Fundamental particles maintain quantum coherence because they are geometrically shallow — their confirmation records are brief relative to the full causal stack depth.

This constitutes a unified geometric account of both the GUP minimum length and quantum decoherence as consequences of a single structural property — causal depth and polarization complexity in the FCC down-vector architecture. The derivation of the explicit functional forms relating β and decoherence time to the joint causal-depth and polarization-complexity profile is reserved for subsequent work.

Conclusion: A Complete Generative Mechanism

The Triadic Measurement Substrate presents a minimal, self-consistent generative mechanism for wave dynamics and binary computational content in three dimensions. Beginning from five axioms and the Principle of Generative Economy, the following results emerge necessarily:

- Axioms 0–2 $\rightarrow$ sphere primitive $\rightarrow$ hexagonal 2D packing $\rightarrow$ triadic confirmation
- Triadic confirmation $\rightarrow$ the 1 $\rightarrow$ 4 multiplier, exact and remainder-free by geometric necessity
- Triadic confirmation pulling from the maximum-entropy substrate $\rightarrow$ binary alphabet $\bar{B} \in \{0, 1\}$ as the irreducible computational content of every confirmed bit, with polarization $\chi \in \{-1, +1\}$ transported by every touch vector signal
- Axiom 4 (causal anchoring) $\rightarrow$ Z-displacement orthogonal by sphere-contact geometry $\rightarrow$ ABCABC stacking $\rightarrow$ FCC lattice

- FCC lattice $\rightarrow$ tetrahedral confirmation units $\rightarrow$ $[[4,2,2]]$ structural correspondence $\rightarrow$ faithful encoding of value (LQ1) and causal-history (LQ2) $\rightarrow$ fault-tolerant propagation of both
- 2D agent manifold $\rightarrow$ minimum resolvable length identified with ℓ_p (Planck length)
- Planck length + speed of light $\rightarrow$ $\tau_{\text{flop}} = t_p$ (Planck time)
- FCC face-diagonal geometry $\rightarrow$ $c_{\text{TMS}} = 1/\sqrt{2}$, forced by lattice structure
- $1 \rightarrow 4$ multiplier $\rightarrow$ self-amplifying cascade regulated by U_{sh} $\rightarrow$ co-generated conjugate pair $(\varphi, \Pi) \rightarrow$ second-order dynamics
- Fault-tolerant FCC lattice + second-order dynamics $\rightarrow$ continuum limit $\rightarrow \partial^2 \varphi / \partial t^2 - c^2 \nabla^2 \varphi = 0$
- Polarization field ψ as parallel object with identical propagation $\rightarrow$ substrate for computational dynamics developed in Volume 1, Paper 2
- Causal depth d and polarization complexity K_ψ in down-vector structure $\rightarrow \beta = f(d, K_\psi) \rightarrow$ GUP and decoherence from a single geometric property

Every quantity appearing in the TMS — the absolute length scale, the absolute time scale, the propagation speed, the binary alphabet — is fixed by the axioms and the Principle of Generative Economy without external input. The confirmation rules and lattice connectivity that govern the system’s behavior are not assumptions imported from outside the framework; they are the geometric consequences of the primitives established in Section I. The model has no numerical degrees of freedom left to tune.

The single fundamental rate of the system — the flop — governs confirmation, synthesis, and propagation simultaneously. Flop time τ is not a background parameter but the counted succession of confirmation events. Space is not a container but the accumulated causal record of prior states. The $[[4,2,2]]$ structural correspondence ensures the system simultaneously maintains the present state, the binary value of every confirmed bit, and the integrity of the causal record in every flop. The wave equation emerges from the geometry of that record propagating forward.

The TMS establishes that a discrete, informationally closed confirmation substrate built from geometric primitives is both necessary and sufficient to generate wave dynamics and a binary computational alphabet from which physical laws may be derived. The paper opens with the incompressibility of U_{sh} (Axiom 0) and closes with a falsifiable prediction — causal-depth and polarization-complexity-dependent GUP deformation — grounded in that same incompressibility. The arc from maximum entropy ground state to Planck-scale uncertainty is a single unbroken geometric chain.

The derivation of further physical laws (electromagnetism, gravity), the dynamics of the polarization field ψ and its coupling to confirmation events, the explicit functional form of $f(d, K_\psi)$, and the scaling properties of the triadic confirmation rule across complex systems are left for subsequent work in Volume 1.

Acknowledgments

The development of this volume was substantially aided by extended collaboration with several large language models used as critical sounding boards, pedagogical resources, formal-rigor checkers, and external working memory across many sessions during 2024–2026.

Gemini (Google DeepMind) was the long-running primary collaborator across the framework’s conceptual development, providing physics pedagogy, structural critique, and ongoing review of the ideas that became Volume 1. The author’s path into the technical literature, and the working-out of the framework’s core conceptual moves over a long period preceding formalization, were carried by extended dialogue with Gemini.

Claude (Anthropic) was the primary collaborator across the formalization, derivation tightening, simulation work, citation auditing, and editorial passes that brought the volume to its present form, including the sublattice bifurcation analysis empirically verified in Paper 5 Appendix E.

ChatGPT (OpenAI), DeepSeek, Grok (xAI), and Granite (IBM) contributed additional structural feedback and review at various points.

All theoretical commitments, the choice of axioms, the Principle of Generative Economy, the sublattice bifurcation insight, and the framework's overall architecture are the author's; the AI systems' role was to teach, formalize, critique, verify, and document.

References

- Aitchison, C. T., & Béri, B. (2025). Spacetime Spins: Statistical mechanics for error correction with stabilizer circuits. arXiv preprint. <https://arxiv.org/abs/2512.21991>
- El Morsalani, M. (2025). Encoding Spacetime: A Comprehensive Review of Holographic Quantum Codes. Preprint.
- Fejes Tóth, L. (1940). Über einen geometrischen Satz. *Mathematische Zeitschrift*, 46(1), 83–85. doi:10.1007/BF01181430.
- Giné, J. (2026). On the Possibility of Quantum Gravity Emerging from Geometry. arXiv preprint. <https://arxiv.org/abs/2602.16219>
- Goldstein, H., Poole, C., & Safko, J. (2002). *Classical Mechanics* (3rd ed.). Addison-Wesley.
- Gottesman, D. (1997). Stabilizer Codes and Quantum Error Correction. Ph.D. Thesis, California Institute of Technology. <https://arxiv.org/abs/quant-ph/9705052>
- Hamilton, W. R. (1834). On a General Method in Dynamics. *Philosophical Transactions of the Royal Society*, 124, 247–308.
- Hatcher, A. (2002). *Algebraic Topology*. Cambridge University Press.
- Hossenfelder, S. (2013). Minimal Length Scale Scenarios for Quantum Gravity. *Living Reviews in Relativity*, 16(1), 2. doi:10.12942/lrr-2013-2.
- Jaynes, E. T. (1957). Information Theory and Statistical Mechanics. *Physical Review*, 106(4), 620–630. doi:10.1103/PhysRev.106.620.
- Kauffman, S. (1993). *The Origins of Order*. Oxford University Press.
- Kolmogorov, A. N. (1965). Three approaches to the quantitative definition of information. *Problems of Information Transmission*, 1(1), 1–7.
- Levin, M. (2019). The Computational Boundary of a Self: Developmental Bioelectricity Drives Multicellularity and Scale-Free Cognition. *Frontiers in Psychology*, 10, 2688. doi:10.3389/fpsyg.2019.02688.
- Ma, Y., Tsao, D., & Shum, H.-Y. (2022). On the Principles of Parsimony and Self-Consistency for the Emergence of Intelligence. *Frontiers of Information Technology & Electronic Engineering*, 23, 1298–1323. <https://arxiv.org/abs/2207.04630>
- Pastawski, F., Yoshida, B., Harlow, D., & Preskill, J. (2015). Holographic quantum error-correcting codes: toy models for the bulk/boundary correspondence. *Journal of High Energy Physics*, 2015(6), 149. doi:10.1007/JHEP06(2015)149. <https://arxiv.org/abs/1503.06237>
- Rissanen, J. (1978). Modeling by shortest data description. *Automatica*, 14(5), 465–471. doi:10.1016/0005-1098(78)90005-5.
- Shannon, C. E. (1948). A mathematical theory of communication. *Bell System Technical Journal*, 27(3), 379–423. doi:10.1002/j.1538-7305.1948.tb01338.x.

Thue, A. (1890). Om nogle geometrisk-taltheoretiske Theoremer. Forhandlingerne ved de Skandinaviske Naturforskere, 14, 352–353.

Wheeler, J. A. (1990). Information, physics, quantum: The search for links. In W. H. Zurek (Ed.), *Complexity, Entropy and the Physics of Information* (pp. 3–28). Addison-Wesley.

Wolfram, S. (2002). *A New Kind of Science*. Wolfram Media.